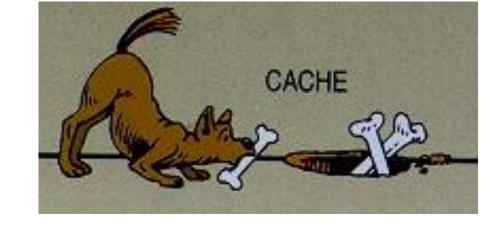
ENCE360 Operating Systems



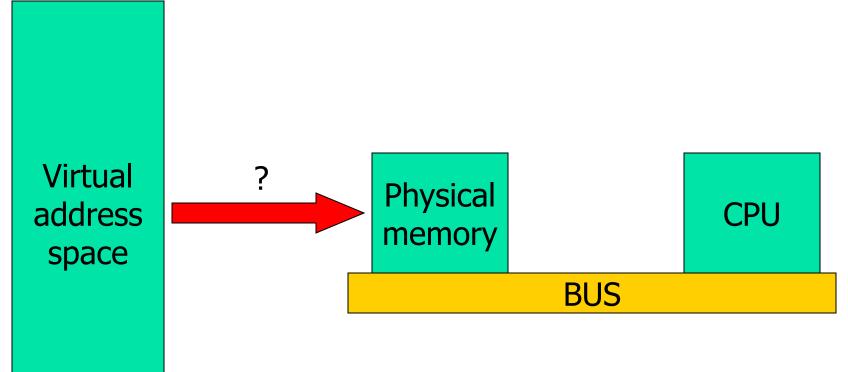


Memory Management

Caching for Faster Memory

MOS Ch 3

Virtual memory



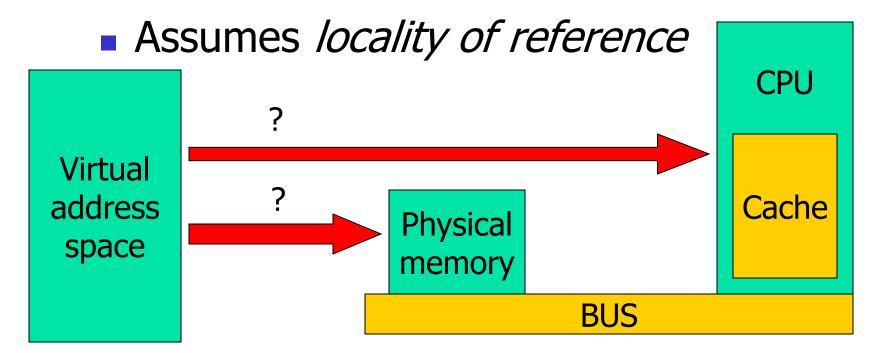
Max 64 bit number = 18, 446,744, 073,709, 551,615
As a memory address, access 18 million Terabytes of RAM
(= 18,000 Petabytes = 18 Exabytes)

Virtual Memory

- Programmer access *virtual* addresses
- Hardware translates into physical addresses
- Bus retrieves the data from physical memory
- Bus transaction moves 1 or more bytes of data in from memory

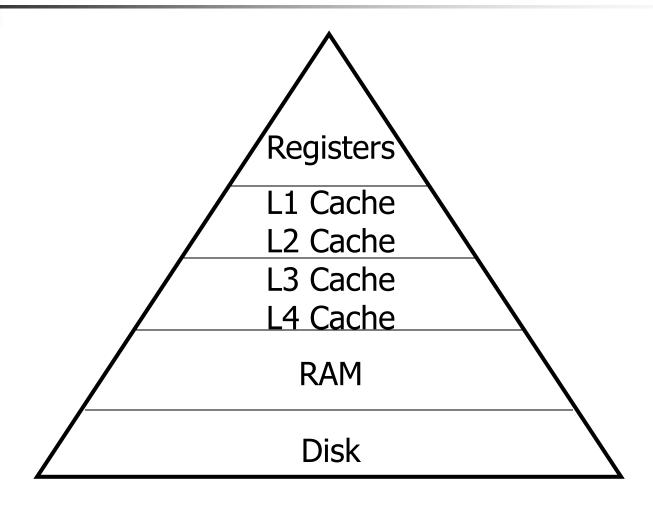
Caching

- Memory much slower than processor
- Put faster "cache" close to CPU



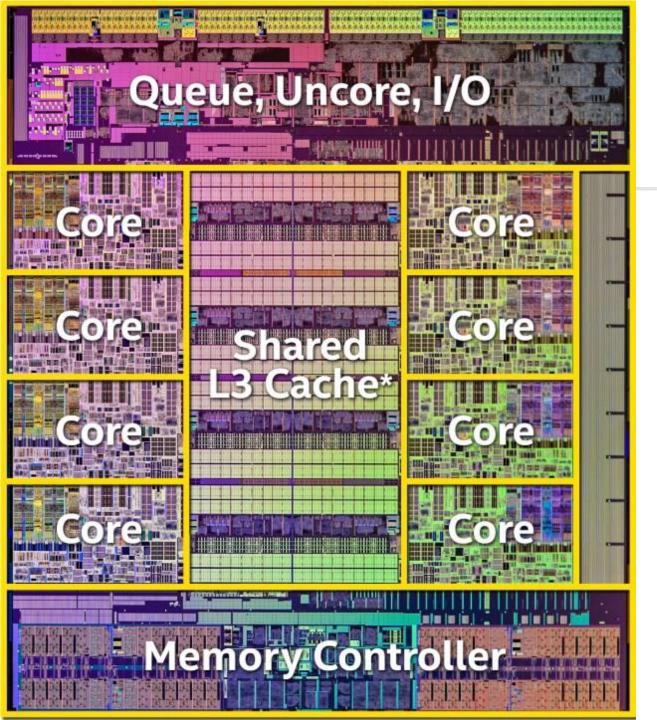


Memory hierarchy



Supporting multiple cores

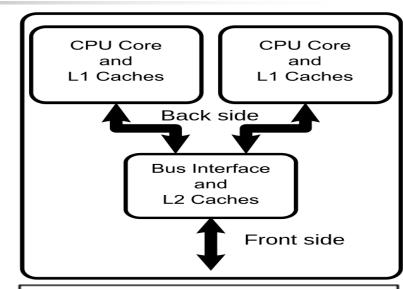
- 4
- past: at its simplest level, an L3 cache was just a larger, slower version of the L2 cache. Back when most chips were singlecore processors, this was generally true - first L3 caches were actually built on the motherboard itself, connected to the CPU via the backside bus...but now:
- L1 cache per core (e.g. L1 32KB, L2 256KB per core, L3 8MB shared)
- 1999 L2 cache moved on-die (per core, but can be per 2 cores)
- 2008 L3 cache shared between all the cores (8MB L3/4 cores)
- 2013 L4 cache exists off-die, but on-package (uncommon)
 L4 acts as a victim cache to the L3 cache, meaning anything evicted from L3 cache immediately goes into the L4 cache

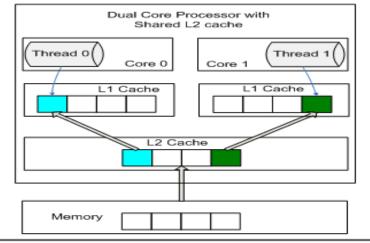


Intel Haslwell

Supporting multiple cores

- Dual-core processor with CPUlocal level-1 caches and a shared, on-die level-2 cache
- The benefits of such a shared cache system include:
 - Flexibility for programmers
 - Reduce cache-coherency complexity
 - Reduce data-storage redundancy
 - Reduce front-side bus traffic





4

Cache entries

Cache line: fixed length (4 to 64 bytes)

Valid bit (covered later)

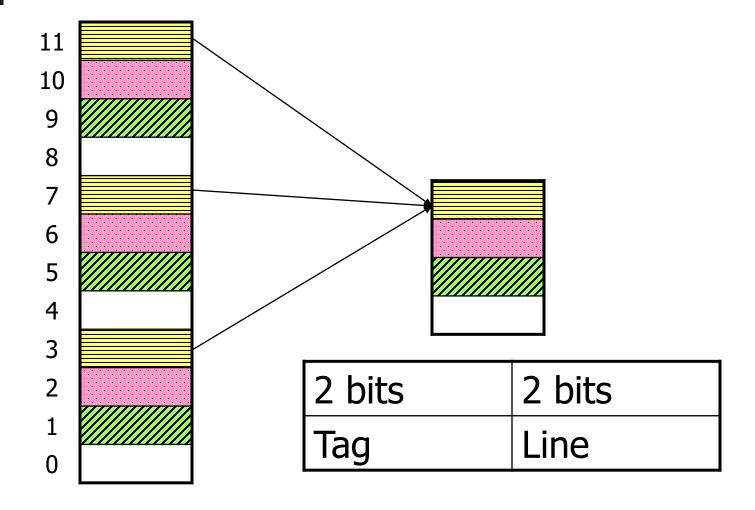
Cache tag: identifies address range

| Valid | Tag | Data 1 | Data 2 | Data 3 | Data 4 |
|-------|-----|--------|--------|--------|--------|
|-------|-----|--------|--------|--------|--------|

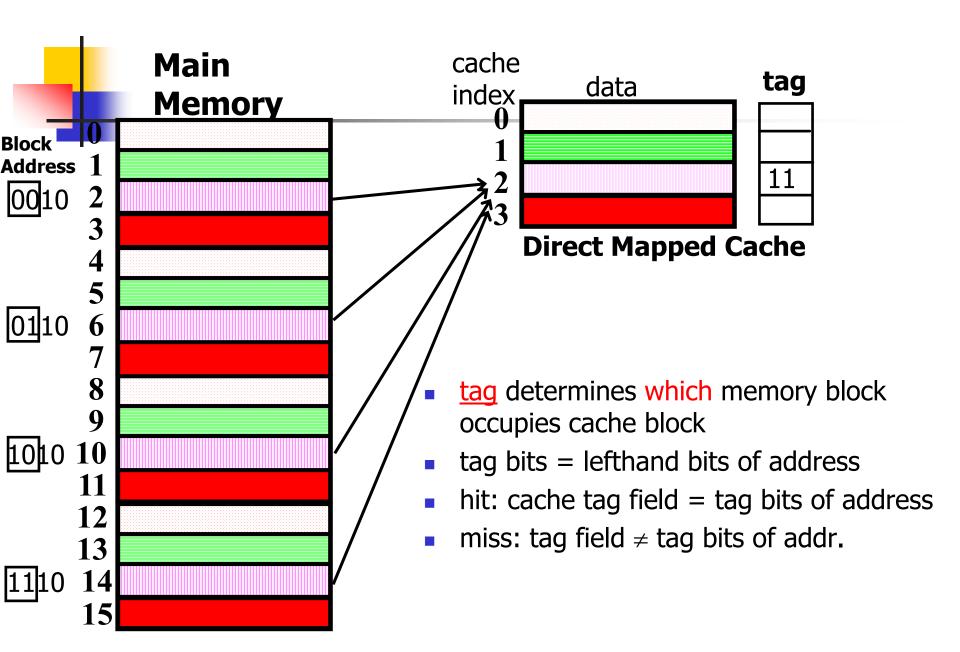
Cache flavours

- Direct mapping
 - Address indicates where it is (fast!)
- Associative
 - Any cache line for any address
 - Address identified by tag
- Set associative
 - Combination of the two





Simplest Cache: Direct Mapped w/Tag





Direct mapping pros and cons

- Pros:
 - FAST! Direct access to the cache
 - Simple
- Cons:
 - thrashing



Associative caches

- Cache entry is "associated" with address by the tag
 - Search all entries, OR
 - Comparator circuit for each cache line

| Tag | Byte |
|--------|------|
| 001101 | 01 |

| Tag | D0 | D1 | D2 | D3 |
|--------|----|----|----|----|
| 000010 | | | | |
| 101001 | | | | |
| 001101 | | | | |
| 001110 | | | | |



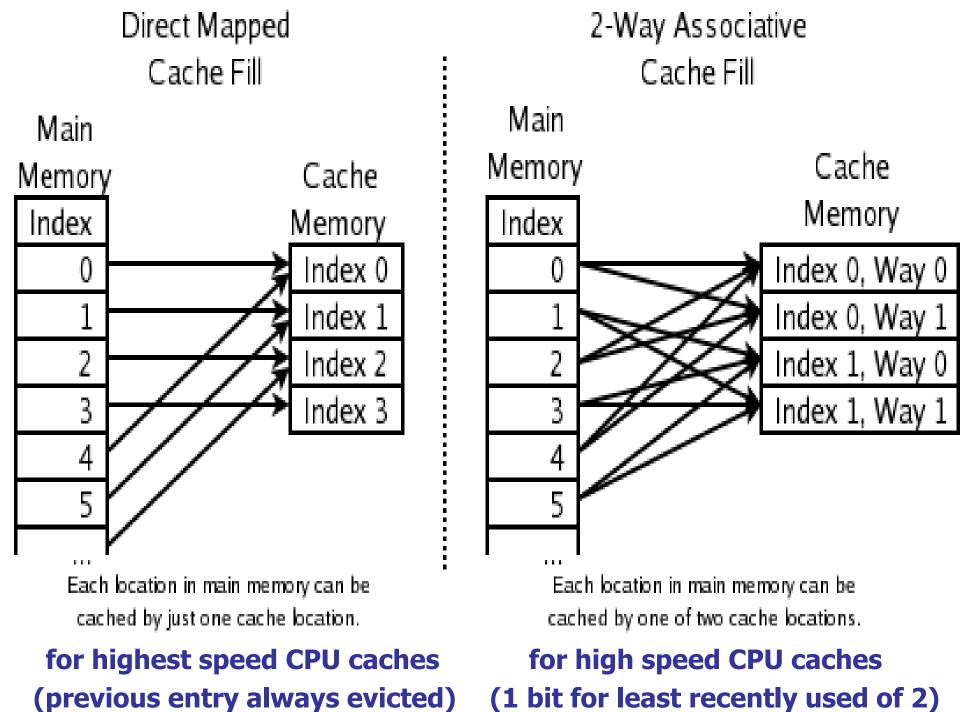
Associative cache

- Pros:
 - Maximises cache use
 - Minimises collisions

- Cons
 - Slow, OR
 - Complex/expensive



- Combination of direct and associative
- Direct-mapped set of associative caches
- N-way cache: n associative entries per direct address map



Set-associative example

| | ٧ | Tag | Data | ٧ | Tag | Data |
|---|---------|------|------|---------|------|------|
| 7 | | 0011 | | | 0101 | |
| 6 | | 1010 | | | 1100 | |
| 5 | | 1101 | | | 0001 | |
| 4 | | 0101 | | | 0010 | |
| 3 | | 1111 | | | 1001 | |
| 2 | | 0000 | | | 0101 | |
| 1 | | 1010 | | | 1111 | |
| 0 | | 1101 | | | 0000 | |
| | Entry A | | | Entry B | | |



Set-associative pros and cons

Pros

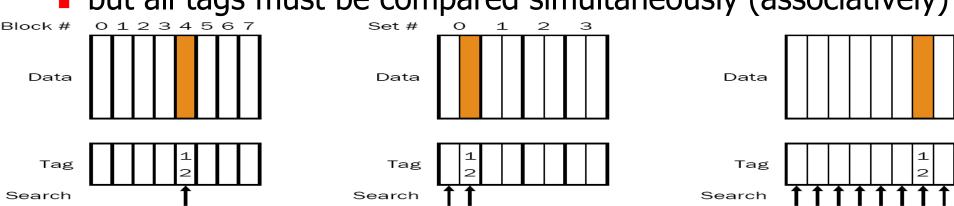
- High hit rate (8-way = 60% lower miss rate)
- Easier to build / faster than fullyassociative

Cons

Slower access (if clocked)

Direct Mapped (D) vs Fully Associative (A) vs Set Associative (SA) Caching

- each address can be stored in one cache location
 - address = memory block (tag) + cache line (+ offset in line)
- 2 way SA: each address stored in 2 cache locations
 - address = memory block (tag) + set number (+ offset in line)
- A: any address can be stored in any cache location
 - address = line in memory (+ offset in line)
 - but all tags must be compared simultaneously (associatively)



Direct/Set/Associate Relationship

Conceptually, the direct mapped and fully associative caches are just "special cases" of the N-way set associative cache.

You can set "N" to 1 to make a "1-way" set associative cache. If you do this, then there is only one line/block per set, which is the same as a direct mapped cache because each memory address is back to pointing to only one possible cache location.

On the other hand, suppose you make "N" really large; say, you set "N" to be equal to the number of entries in the cache. If you do this, then you only have one set, containing all of the lines/blocks, and every memory location points to that huge set. This means that any memory address can be in that one (read any) cache entry, and you are back to a fully associative cache.

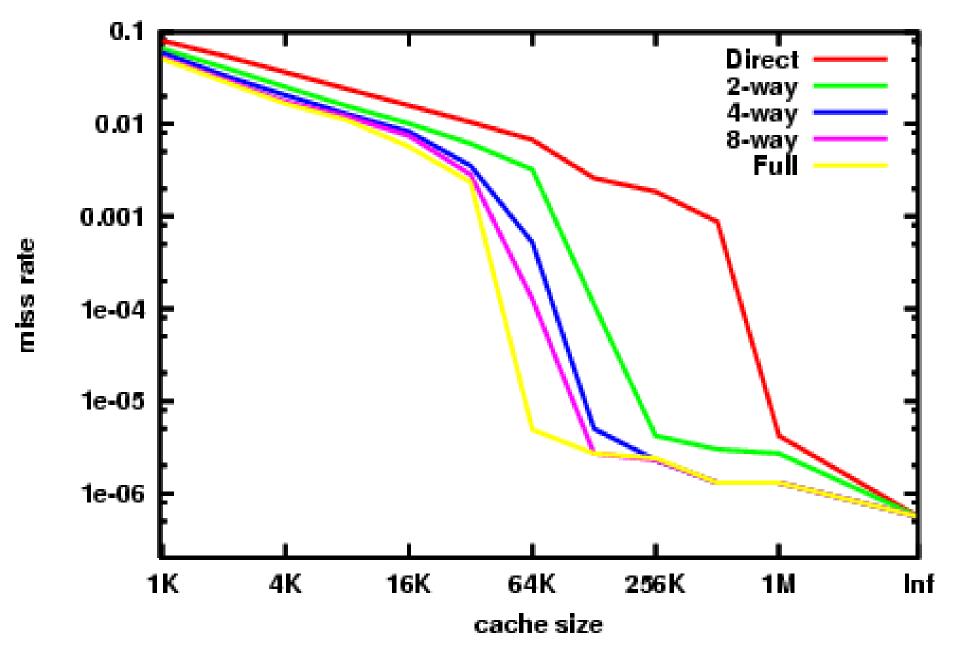
Cache Performance

| Cache Type | Hit Ratio | Search Speed |
|-------------------------------|-------------------------------------|-------------------------------|
| Direct Mapped | Good | Best |
| Fully Associative | Best | Moderate |
| N-Way Set Associative, N>1 | Very Good, Better as N Increases | Good, Worse as N Increases |

Direct mapped and set associative caches are by far the most common.

Direct mapping is used more for level 3 or 4 caches on motherboards, while the higher-performance set-associative cache is found more commonly on the smaller primary caches contained within processors.

Also visit: www.pcguide.com/ref/mbsys/cache/func.htm

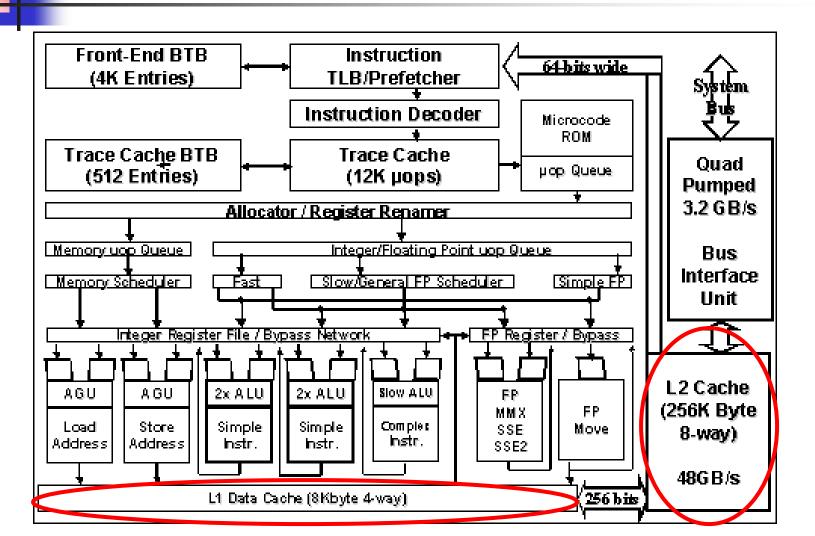


Miss rate versus cache size on the Integer portion of SPEC CPU2000 A clever compiler can take cache collisions into account, placing instructions and data in memory.

Cache architectures

- Multi-level:
 - L1 on-die (private to each core)
 - L2 on-die (private to each core)
 - L3 on-die (shared between cores)
 - L4 off-die, but on-package (caches L3 evictions)
- Unified vs split
 - Unified: better usage (% full)
 - Split: different types/policies

Example: Pentium (single core)





Cache Position



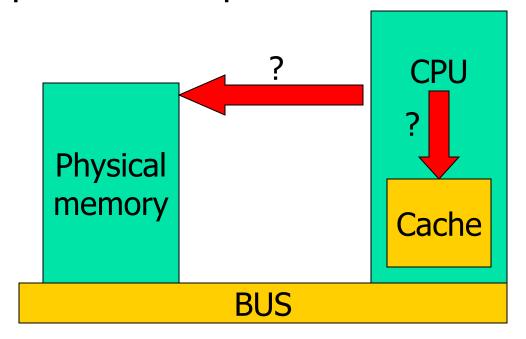
- Physical vs Virtual:
 - Physical:
 - address translated first
 - Context not required
 - Virtual:
 - In parallel with address translation
 - But: context required

Replacement Policy

- Which cache entry to "evict"?
 - Direct mapping: no issue
 - Associative/set associative:
 - RANDOM
 - FIFO
 - Least recently used (LRU)
 - Least frequently used

And then it all went horribly wrong...

- Policy needed for writing to memory
 - Consistency cache and memory "agree"
 - Speed multiple writes



Write policies

Write-through

always write to disk

Copy back or Write-back

 write to disk only when evicting (even if unchanged)

Write-deferred

- Mark cache entry as "dirty" (if changed)
- Write out dirty entry on eviction

Write allocation

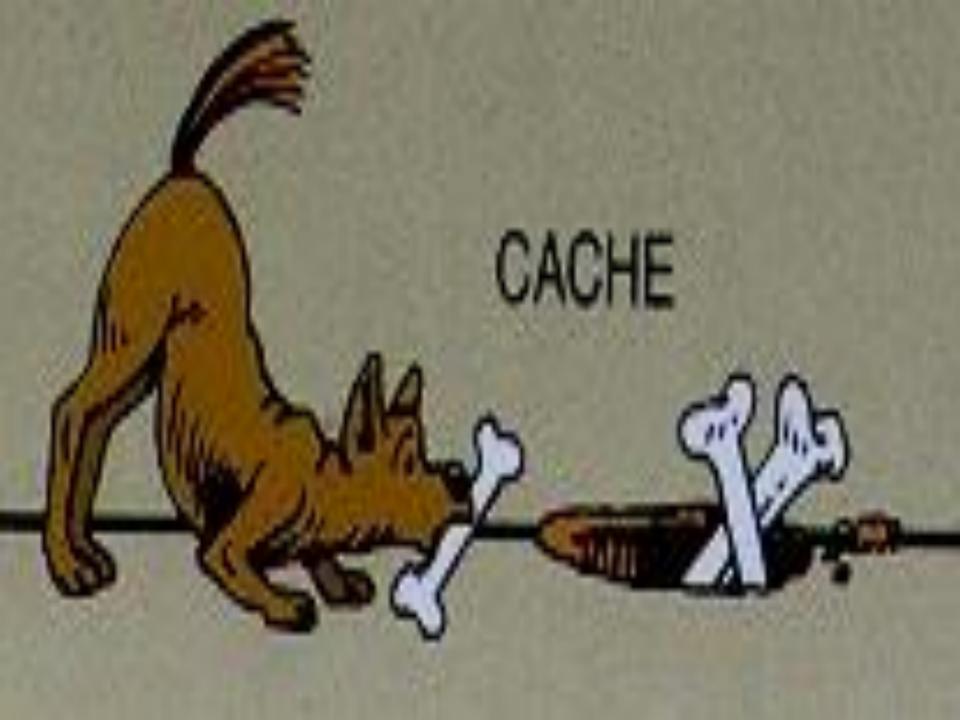
bring missed entry into cache

Prefetching

- Load next n lines of memory into cache
- E.g. video
- May write over needed data
- Memory traffic increases
- Processor idle during fetch?
- Clever hardware can make worthwhile

Summary

- Most modern processors/operating systems use caches, usually a least two levels
- Cache design is complex
- Design affects performance
- Need to know cache design to optimise code



Example Exam Question

The following is a fragment of a **16KByte**, **2-way associative L1 cache with a line size of 16 bytes**, an update policy of write-deferred, write-allocate, and a replacement policy of least-recently-used (LRU).

| Set | Tag | Valid? | Dirty? | Tag | Valid? | Dirty? |
|-----------|-----|--------|---------------|-----|--------|--------|
| 47 | 101 | Y | N | 111 | N | N |
| 46 | 010 | Y | N | 110 | Y | Y |
| 45 | 011 | Y | Y | 000 | N | N |
| 44 | 111 | Y | Y | 101 | Y | N |
| 43 | 010 | Y | N | 111 | Υ | N |

For each of the following consecutive memory operations (in binary) for a 64KByte virtual memory space, indicate the *number of memory transfers* assuming that the bus width is 4 bytes, and show the **state of the "valid"** and "dirty" bits of the corresponding cache entry after the operation. *Show your working.*

| (a) [2 marks] | WRITE: | 1010001011110100 |
|----------------------|--------------|------------------|
| (b) [2 marks] | READ: | 1100001011100111 |
| (c) [2 marks] | READ: | 1100001011110000 |
| (d) [2 marks] | WRITE: | 0000001011111111 |
| (e) [2 marks] | READ: | 0100001011011010 |

Example Exam Question

The following is a fragment of a **16KByte**, **2-way associative L1 cache with a line size of 16 bytes**, an update policy of write-deferred, write-allocate, and a replacement policy of least-recently-used (LRU).

| Set | Tag | Valid? | Dirty? | | Tag | Valid? | Dirty? | |
|-----------|------------|--------|--------|-------------|-----|--------|---------------|------|
| 47 | 101 | Y | N | 16KB | 111 | N | N | 16KB |
| 46 | 010 | Y | N | 16KB | 110 | Y | Y | 16KB |
| 45 | 011 | Y | Y | 16KB | 000 | N | N | 16KB |
| 44 | 111 | Y | Y | 16KB | 101 | Y | N | 16KB |
| 43 | 010 | Y | N | 16KB | 111 | Y | N | 16KB |

For each of the following consecutive memory operations (in binary) for a 64KByte virtual memory space, indicate the *number of memory transfers* assuming that the bus width is 4 bytes, and show the **state of the "valid"** and "dirty" bits of the corresponding cache entry after the operation. *Show your working.*

| (a) [2 marks] | WRITE: | 1010001011110100 |
|----------------------|--------------|------------------|
| (b) [2 marks] | READ: | 1100001011100111 |
| (c) [2 marks] | READ: | 1100001011110000 |
| (d) [2 marks] | WRITE: | 0000001011111111 |
| (e) [2 marks] | READ: | 0100001011011010 |



...16KByte, 2-way associative L1 cache with a line size of 16 bytes...

- 16 byte line/block size 4 bits to address each byte
- 2-way associative means 32 bytes per set (2 x 16)
- 16KB cache @ 32 bytes per set = 512 sets (9 bits) $(16KB/32B = 2^{14}/2^5 = 2^9 \text{ sets} = 512 \text{ sets})$
- 9 bits to represent set number leaves 3 bits for tag
- low 4 bits = offset (to address each byte in a line/block)
- next 9 bits = set number (cache entry number)
- high 3 bits = tag (block number in memory)

```
3 bits 9 bits 4 bits tag set num offset (a) WRITE 101 000101111 0100
```

Example Exam Question

The following is a fragment of a **16KByte**, **2-way associative L1 cache with a line size of 16 bytes**, an update policy of write-deferred, write-allocate, and a replacement policy of least-recently-used (LRU).

| Set | Tag | Valid? | Dirty? | | Tag | Valid? | Dirty? | |
|-------|-----|--------|--------|-------------|-----|--------|--------|-------------|
| 511 | | | • | | | | | |
| l. I | | | | | | | | |
| - | | | | | | | | |
| l. I | | | | | | | | |
| 47 | 101 | Y | N | 16KB | 111 | N | N | 16KB |
| 46 | 010 | Y | N | 16KB | 110 | Y | Y | 16KB |
| 45 | 011 | Y | Y | 16KB | 000 | N | N | 16KB |
| 44 | 111 | Y | Y | 16KB | 101 | Y | N | 16KB |
| 43 | 010 | Y | N | 16KB | 111 | Y | N | 16KB |
| I. I | | | | | | | | |
| I. I. | | | | | | | | |
| I. I | | | | | | | | |
| | | | | | | | | |

... for 512 sets ...

```
3 bit 9 bit 4 bit tag set num offset (a) WRITE 101 000101111 0100
```

- Set number 1011111_2 = set 47_{10} (decimal)
- Tag = 101 = Tag matched = Cache hit that line/block now becomes dirty:

| Set | Tag | Valid? | Dirty? | Tag | Valid? | Dirty? |
|-----------|-----|--------|--------|-----|--------|--------|
| 47 | 101 | Y | N | 111 | N | N |
| 46 | 010 | Υ | Ν | 110 | Υ | Υ |
| 45 | 011 | Υ | Υ | 000 | N | N |
| 44 | 111 | Υ | Υ | 101 | Υ | N |
| 43 | 010 | Υ | Ν | 111 | Υ | Ν |

- No memory transfers because not write through.
- Valid = Y
- Dirty = Y

(b) READ: 110 000101110 0100

- Set number 1011110_2 = set 46_{10}
- Tag = 110 = Tag matched = Cache hit:

| Set | Tag | Valid? | Dirty? | Tag | Valid? | Dirty? |
|-----------|-----|--------|--------|------------|--------|--------|
| 47 | 101 | Υ | Υ | 111 | N | N |
| 46 | 010 | Υ | N | 110 | Y | Y |
| 45 | 011 | Υ | Υ | 000 | N | N |
| 44 | 111 | Υ | Υ | 101 | Υ | N |
| 43 | 010 | Υ | N | 111 | Υ | N |

- No memory transfers required cache remains the same
- Valid = Y
- Dirty = Y

(c) READ 110 000101111 0000

- Set number 1011111_2 = set 47_{10}
- Tag = 110 = No tag matched = Cache miss as no tag 110 is found on set 47 so need to bring one line/block into the cache to replace an invalid (empty) entry (or a least recently used entry):

| Set | Tag | Valid? | Dirty? | Tag | Valid? | Dirty? |
|-----------|-----|--------|--------|-----|--------|--------|
| 47 | 101 | Υ | Υ | 111 | N | N |
| 46 | 010 | Υ | N | 110 | Υ | Υ |
| 45 | 011 | Υ | Υ | 000 | N | N |
| 44 | 111 | Υ | Υ | 101 | Υ | N |
| 43 | 010 | Υ | N | 111 | Υ | N |

- 4 memory transfers because must read in 16 bytes over 32-bit bus
- Valid = Y
- Dirty = N

(d) WRITE 000 000101111 1111

- Set number 1011111₂ = set 47₁₀
- Tag = 000 = No tag matched = Cache miss as no tag 000 is found on set 47 – so need to bring one line/block into the cache to replace the least recently used entry:

| Set | Tag | Valid? | Dirty? | Tag | Valid? | Dirty? |
|-----------|-----|--------|--------|-----|--------|--------|
| 47 | 101 | Y | Y | 110 | Υ | N |
| 46 | 010 | Υ | N | 110 | Υ | Υ |
| 45 | 011 | Υ | Υ | 000 | Ν | N |
| 44 | 111 | Υ | Υ | 101 | Υ | N |
| 43 | 010 | Υ | N | 111 | Υ | N |

- 8 memory transfers: dirty & valid so write it out 4 memory transfers and then bring in the new line/block – 4 memory transfers
- Valid = Y
- Dirty = Y

(e) READ 010 000101101 1010

- Set number 101101₂ = set 45₁₀
- Tag = 010 = No tag matched = Cache miss as no tag 010 is found on set 45 – so need to bring one line/block into the cache to replace an invalid (empty) entry or a least recently used entry:

| Set | Tag | Valid? | Dirty? | Tag | Valid? | Dirty? |
|-----------|-----|--------|--------|-----|--------|--------|
| 47 | 000 | Υ | Υ | 110 | Υ | N |
| 46 | 010 | Υ | N | 110 | Υ | Υ |
| 45 | 011 | Υ | Υ | 000 | N | N |
| 44 | 111 | Υ | Υ | 101 | Υ | N |
| 43 | 010 | Υ | N | 111 | Υ | N |

- 4 memory transfers because must read in 16 bytes over 32-bit bus
- Valid = Y
- Dirty = N

number of memory bus transfers

- (a) Write set 47, tag 101 **0**
- (b) Read set 46, tag 110 **0**
- (c) Read set 47, tag 110 **4**
- (d) Write set 47, tag 000 **16**
- (d) Read set 45, tag 010 **4**

| Set | Tag | Valid? | Dirty? | Tag | Valid? | Dirty? |
|-----|-----|--------|--------|-----|--------|--------|
| 47 | 000 | Υ | Y | 110 | Y | N |
| 46 | 010 | Υ | N | 110 | Υ | Υ |
| 45 | 011 | Υ | Υ | 010 | Y | N |
| 44 | 111 | Υ | Υ | 101 | Υ | Ν |
| 43 | 010 | Υ | N | 111 | Υ | N |

Example Exam Question

Explain what each of the following mean, and give one example for each of how they are taken advantage of in a modern computer system (note: it is not sufficient to name a processor/operating system that uses it):

- (a) [2 marks] Delegation
- (b) [2 marks] Locality of reference



Solution: Delegation

- The CPU/processor offloads (delegates) work to co-processors.
- Examples: Disk controller, FPU, video/graphics card



Solution: Locality of reference

- Memory accesses (data or instructions) are likely to be near (spatially and temporally).
- Example: Caches take advantage of this to speed up memory access by paging in surrounding memory.

